

Obed's stock market analysis using weather

April 4, 2024

```
[1]: %%capture
      %pip install tqdm seaborn skillsnetwork scikit-learn==0.24
```

```
[2]: from functools import reduce
      from copy import deepcopy
      import tqdm
      import numpy as np
      from scipy.signal import periodogram
      from scipy.stats import binomtest
      import pandas as pd
      import skillsnetwork
      import statsmodels.api as sm
      from statsmodels.tsa.stattools import adfuller, kpss
      from sklearn.model_selection import TimeSeriesSplit
      from sklearn.metrics import mean_absolute_error
      from sklearn.ensemble import RandomForestRegressor
      import matplotlib.pyplot as plt
      from matplotlib.patches import Patch
      import seaborn as sns
      %matplotlib inline

      # Float format for pandas display
      pd.set_option('display.float_format', lambda x: '%.8f' % x)

      # Suppress unneeded warnings:
      def warn(*args, **kwargs):
          pass
      import warnings
      warnings.warn = warn
      warnings.filterwarnings('ignore')

      sns.set_context('notebook')
      sns.set(style="darkgrid")
```

```
[3]: # Import weather data
      # Note that all of the columns are imported as strings
```

```

# This is generally the safest option, but requires additional processing later
↳ on

await skillsnetwork.download_dataset(
    'https://cf-courses-data.s3.us.cloud-object-storage.appdomain.cloud/
    ↳ IBMSkillsNetwork-GPXXOK1YEN/laguardia.csv'
)
laguardia = pd.read_csv('laguardia.csv', dtype='str')

# Import DOW Jones Industrial Average historical data

await skillsnetwork.download_dataset(
    'https://cf-courses-data.s3.us.cloud-object-storage.appdomain.cloud/
    ↳ IBMSkillsNetwork-GPXXOK1YEN/dow_jones.csv'
)
dow = pd.read_csv('dow_jones.csv', dtype='str')

```

```

Downloading laguardia.csv: 0%|          | 0/34039379 [00:00<?, ?it/s]
Saved as 'laguardia.csv'
Downloading dow_jones.csv: 0%|          | 0/1635081 [00:00<?, ?it/s]
Saved as 'dow_jones.csv'

```

```

[4]: # Weather data
laguardia['DATE'] = pd.to_datetime(laguardia.DATE)
laguardia[['wind',
            'dew_point',
            'temp', 'pressure',
            'cloud_cover']] = laguardia[['wind',
                                            'dew_point',
                                            'temp',
                                            'pressure',
                                            'cloud_cover']].astype(float)

# Market data
dow['DATE'] = pd.to_datetime(dow.DATE)
# Drop missing value rows
dow = dow.loc[dow.Open != 'na']
dow[[i for i in dow.columns if i != 'DATE']] = dow[[i for i in dow.columns if i
    ↳ != 'DATE']].astype(float)
dow['Volume'] = dow.Volume.astype(int)

[5]: laguardia = laguardia.loc[:, ['DATE', 'temp', 'cloud_cover']]
dow = dow.loc[:, ['DATE', 'Close']]

```

```
[6]: # Print the `DATE` field in the `laguardia` dataset:
print("laguardia 'DATE' field head")
print(laguardia.DATE.head())

# The following code shows the hours for which data is available
print("\n laguardia 'DATE' field hour availability")
print(sorted(laguardia.DATE.dt.hour.unique()))

# The following code shows the minutes for which data is available
print("\n laguardia 'DATE' field minute availability")
print(sorted(laguardia.DATE.dt.minute.unique()))
```

laguardia 'DATE' field head

```
0    1948-07-01 11:00:00
1    1948-07-01 12:00:00
2    1948-07-01 13:00:00
3    1948-07-01 14:00:00
4    1948-07-01 15:00:00
```

Name: DATE, dtype: datetime64[ns]

laguardia 'DATE' field hour availability

```
[0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, 18, 19, 20, 21,
22, 23]
```

laguardia 'DATE' field minute availability

```
[0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, 18, 19, 20, 21,
22, 23, 24, 25, 26, 27, 28, 29, 30, 31, 32, 33, 34, 35, 36, 37, 38, 39, 40, 41,
42, 43, 44, 45, 46, 47, 48, 49, 50, 51, 52, 53, 54, 55, 56, 57, 58, 59]
```

```
[7]: # Print the `DATE` field in the `dow` dataset:
print("dow 'DATE' field head")
print(dow.DATE.head())

# The following code shows the hours for which data is available
print("\n dow 'DATE' field hour availability")
print(sorted(dow.DATE.dt.hour.unique()))

# The following code shows the minutes for which data is available
print("\n dow 'DATE' field minute availability")
print(sorted(dow.DATE.dt.minute.unique()))
```

dow 'DATE' field head

```
0    2012-12-31
1    2012-12-28
2    2012-12-27
3    2012-12-26
5    2012-12-24
```

Name: DATE, dtype: datetime64[ns]

```
dow 'DATE' field hour availability
[0]
```

```
dow 'DATE' field minute availability
[0]
```

```
[8]: # The following code shows the frequency counts for minutes in `laguardia`:
print("\n laguardia 'DATE' field minute frequency (head):")
print(laguardia.DATE.dt.minute.value_counts().head())
```

```
laguardia 'DATE' field minute frequency (head):
0      503480
51     235096
59      10846
49       3229
30       1714
Name: DATE, dtype: int64
```

```
[9]: print("'laguardia' duplicated:")
print(laguardia.DATE.duplicated().value_counts())
```

```
'laguardia' duplicated:
False      799794
True         2145
Name: DATE, dtype: int64
```

```
[10]: print("'dow' duplicated:")
print(dow.DATE.duplicated().value_counts())
```

```
'dow' duplicated:
False      20962
Name: DATE, dtype: int64
```

```
[11]: print("'laguardia' missing:")
print(laguardia.isna().max())
```

```
'laguardia' missing:
DATE          False
temp           True
cloud_cover    True
dtype: bool
```

```
[12]: pd.set_option('display.float_format', lambda x: '%.2f' % x)
print("'laguardia' description:")
print(laguardia.describe())
```

```

'laguardia' description:
      temp  cloud_cover
count 717298.00    646597.00
mean    13.04         0.60
std     9.76         0.39
min    -19.40         0.00
25%     5.60         0.25
50%    13.30         0.75
75%    21.10         1.00
max     39.40         1.00

```

```

[13]: # The following resamples all data to an hourly frequency by
# taking an average of all minutes that round to that hour.
laguardia['DATE'] = laguardia['DATE'].dt.round('60min')

# Note that a loop is used to account for the fact that each column contains a
# unique set of missing values:
laguardia_cols = []

for c in laguardia.columns:
    if c == 'DATE':
        continue
    else:
        laguardia_cols.append(
            laguardia[['DATE', c]].dropna().groupby(
                'DATE', as_index=False
            ).agg({c: 'mean'})
        )

# Finally, merge all columns back together again:
laguardia_merged = reduce(
    lambda left, right: pd.merge(left, right, on=['DATE'], how='outer'),
    laguardia_cols
)

# Sort by DATE
laguardia_merged.sort_values('DATE', inplace=True)

# Let's see what the merged data looks like:
laguardia_merged.head()

```

```

[13]:          DATE  temp  cloud_cover
0  1948-07-01 11:00:00  22.80         1.00
1  1948-07-01 12:00:00  23.30         0.88
2  1948-07-01 13:00:00  25.00         0.25
3  1948-07-01 14:00:00  26.70         0.00
4  1948-07-01 15:00:00  27.80         0.00

```

```
[14]: laguardia_merged.isna().value_counts()
```

```
[14]: DATE    temp    cloud_cover
      False  False  False          569290
           True    True          44601
           True  False          184
      dtype: int64
```

```
[15]: laguardia_merged[['cloud_cover', 'DATE']].dropna().DATE.diff().value_counts()
```

```
[15]: 0 days 01:00:00    521387
      0 days 03:00:00    41070
      0 days 02:00:00    6354
      0 days 06:00:00     405
      0 days 04:00:00     83
      0 days 05:00:00     44
      0 days 12:00:00     41
      0 days 09:00:00     34
      0 days 07:00:00     15
      0 days 08:00:00     10
      0 days 18:00:00      7
      1 days 00:00:00      5
      0 days 11:00:00      4
      0 days 10:00:00      4
      0 days 13:00:00      2
      0 days 15:00:00      1
      0 days 19:00:00      1
      1 days 01:00:00      1
      0 days 17:00:00      1
      0 days 16:00:00      1
      1 days 12:00:00      1
      0 days 23:00:00      1
      2 days 09:00:00      1
      Name: DATE, dtype: int64
```

```
[16]: laguardia_nan_cloud_cover = laguardia_merged.set_index(
      'DATE', drop=True
    ).sort_index()
laguardia_nan_cloud_cover = laguardia_nan_cloud_cover.reindex(
    pd.date_range(
        start=laguardia_merged.DATE.min(),
        end=laguardia_merged.DATE.max(),
        freq='1H'
    )
)
laguardia_nan_cloud_cover = laguardia_nan_cloud_cover.loc[
    laguardia_nan_cloud_cover.cloud_cover.isna()
]
```

```
]
laguardia_nan_cloud_cover['datetime'] = laguardia_nan_cloud_cover.index
laguardia_nan_cloud_cover.datetime.dt.hour.value_counts()
```

```
[16]: 19    5871
      20    5859
      17    5793
      22    5766
      23    5746
      16    5701
      14    5650
      13    5607
       1    5598
      11    5543
       2    5508
      10    5497
       4    5422
       7    5420
       8    5420
       5    5416
      12     477
      18     458
       0     350
       6     334
       9     258
      21     238
      15     234
       3     198
      Name: datetime, dtype: int64
```

```
[17]: # This should output just one row if there are no missing hours:
      print(laguardia_merged.DATE.diff().value_counts())
```

```
0 days 01:00:00    589778
0 days 03:00:00    23445
0 days 02:00:00      850
1 days 00:00:00        1
      Name: DATE, dtype: int64
```

```
[18]: # Reindex the dataset to remove missing hours
      # First, set the `DATE` column as the index:
      laguardia_merged.set_index('DATE', drop=True, inplace=True)
      # Now reindex
      laguardia_merged = laguardia_merged.reindex(
          pd.date_range(
              start=laguardia_merged.index.min(),
              end=laguardia_merged.index.max(),
```

```

        freq='1H'
    )
)
# Set all data types to float:
laguardia_merged = laguardia_merged.astype(float)

# Interpolate
laguardia_merged.interpolate(method='linear', inplace=True)
laguardia_merged.describe()

```

```

[18]:      temp  cloud_cover
count 661838.00    661838.00
mean      12.92         0.60
std        9.84         0.38
min       -19.40         0.00
25%         5.00         0.25
50%        13.30         0.75
75%        21.10         1.00
max        39.40         1.00

```

```

[19]: laguardia_merged.isna().value_counts()

```

```

[19]: temp  cloud_cover
False  False          661838
dtype: int64

```

```

[20]: # Get weather variables between 8am and 9pm
laguardia_merged_avg = laguardia_merged.between_time('8:00', '9:00').
    ↪reset_index()
laguardia_merged_avg.rename({'index': 'DATE'}, axis=1, inplace=True)
laguardia_merged_avg['DATE'] = laguardia_merged_avg['DATE'].dt.round('1D')
laguardia_merged_avg = laguardia_merged_avg.groupby(
    'DATE', as_index=False
).agg({'temp': 'mean', 'cloud_cover': 'mean'}).set_index('DATE')
rename_dict = dict(
    zip(
        laguardia_merged_avg.columns.tolist(),
        [i + '_avg' for i in laguardia_merged_avg.columns]
    )
)
laguardia_merged_avg.rename(rename_dict, axis=1, inplace=True)
df_weather_final = laguardia_merged_avg
df_weather_final.head()

```

```

[20]:      temp_avg  cloud_cover_avg
DATE
1948-07-02      18.30             0.00

```


1948-07-03	19.40	0.69
1948-07-04	19.40	0.00
1948-07-05	22.50	0.06
1948-07-06	23.05	1.00

```
[21]: # `dow` dataset, gaps between dates (head)
dow.DATE.sort_values().diff().value_counts().head()
```


```
[21]: 1 days    16279
      3 days    3880
      4 days     464
      2 days    325
      5 days      8
      Name: DATE, dtype: int64
```

```
[22]: dow.sort_values('DATE', inplace=True)
df = dow.merge(df_weather_final,
               how='outer',
               left_on='DATE',
               right_index=True).set_index('DATE').sort_index()
df = df.loc[df.index >= df_weather_final.index[0]]
df.sort_index(inplace=True)
df.head()
```


```
[22]:
```

	Close	temp_avg	cloud_cover_avg
DATE			
1948-07-02	190.06	18.30	0.00
1948-07-03	NaN	19.40	0.69
1948-07-04	NaN	19.40	0.00
1948-07-05	NaN	22.50	0.06
1948-07-06	190.55	23.05	1.00


```
[23]: _ = sns.lineplot(data=df.Close).set_title('DJI Close Price')
```



```
[24]: df['log_Close'] = np.log(df.loc[:, 'Close'])  
_ = sns.lineplot(data=df.log_Close).set_title('Log DJI Close Price')
```



```
[25]: log_Close = deepcopy(df.loc[:, 'log_Close'])
log_Close.dropna(inplace=True)
ld_Close = log_Close.diff()
df = df.merge(
    pd.DataFrame(ld_Close).rename({'log_Close': 'ld_Close'}, axis=1),
    how='left',
    left_index=True,
    right_index=True
)
_ = sns.lineplot(data=df.ld_Close).set_title('Log differenced DJI Close Price')
```



```
[26]: print('p-value of ADF test:')
      print(adfuller(df.ld_Close.dropna())[1])
      print('p-value of KPSS test:')
      print(kpss(df.ld_Close.dropna())[1])
```

```
p-value of ADF test:
0.0
p-value of KPSS test:
0.1
```

```
[27]: def plot_periodogram(ts, detrend='linear', ax=None):
      fs = pd.Timedelta("365D6H") / pd.Timedelta("1D")
      frequencies, spectrum = periodogram(
          ts,
          fs=fs,
          detrend=detrend,
          window="boxcar",
          scaling='spectrum',
      )
      if ax is None:
          _, ax = plt.subplots()
```

```

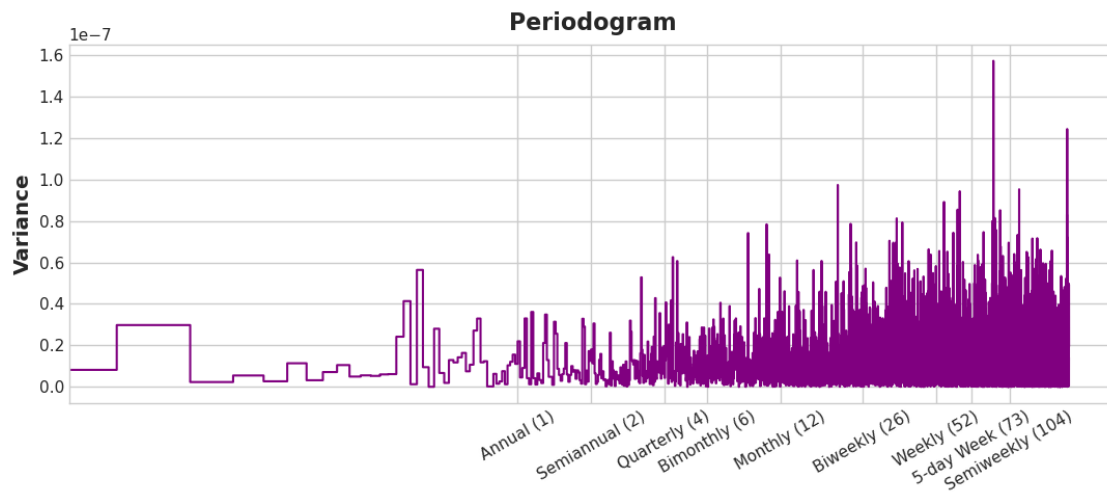
ax.step(frequencies, spectrum, color="purple")
ax.set_xscale("log")
ax.set_xticks([1, 2, 4, 6, 12, 26, 52, 73, 104])
ax.set_xticklabels(
    [
        "Annual (1)",
        "Semiannual (2)",
        "Quarterly (4)",
        "Bimonthly (6)",
        "Monthly (12)",
        "Biweekly (26)",
        "Weekly (52)",
        "5-day Week (73)",
        "Semiweekly (104)",
    ],
    rotation=30,
)
ax.ticklabel_format(axis="y", style="sci", scilimits=(0, 0))
ax.set_ylabel("Variance")
ax.set_title("Periodogram")
return ax

# Set Matplotlib defaults
plt.style.use("seaborn-whitegrid")
plt.rc("figure", autolayout=True, figsize=(11, 5))
plt.rc(
    "axes",
    labelweight="bold",
    labelsiz="large",
    titleweight="bold",
    titlesiz=16,
    titlepad=10,
)
plot_params = dict(
    color="0.75",
    style=".-",
    markeredgecolor="0.25",
    markerfacecolor="0.25",
    legend=False,
)

plot_periodogram(df.loc[:, 'ld_Close'].dropna())


```

[27]: <AxesSubplot:title={'center': 'Periodogram'}, ylabel='Variance'>




```
[29]: plot_periodogram(df.loc[:, 'log_Close'].dropna())
```

```
[29]: <AxesSubplot:title={'center':'Periodogram'}, ylabel='Variance'>
```




```
[30]: _ = sns.lineplot(data=df['temp_avg']).set_title('Average temperature between 8_
         and 9 am')
```



```
[31]: plot_periodogram(df.loc[:, 'temp_avg'].dropna())
```

```
[31]: <AxesSubplot:title={'center':'Periodogram'}, ylabel='Variance'>
```



```
[32]: # Seasonally adjust average temp
y = df.loc[df.index < '1964-01-02', 'temp_avg']
X = [i % 365.25 for i in range(0, len(y.to_numpy()))]
X_full = [i % 365.25 for i in range(0, len(df.temp_avg.to_numpy()))]
degree = 4
coef = np.polyfit(X, y.to_numpy(), degree)
print('Coefficients: %s' % coef)
# create seasonal component
temp_sc_avg = list()
```

```

for i in range(len(X_full)):
    value = coef[-1]
    for d in range(degree):
        value += X_full[i]**(degree-d) * coef[d]
    temp_sc_avg.append(value)


df['temp_sc_avg'] = temp_sc_avg
df['temp_sa'] = df['temp_avg'] - df['temp_sc_avg']

```

Coefficients: [-2.27039916e-08 1.89119633e-05 -4.47566270e-03 2.22570025e-01
1.91433822e+01]

```
[33]: plot_periodogram(df.loc[df.index >= '1964-01-02', 'temp_sa'].dropna())
```

```
[33]: <AxesSubplot:title={'center':'Periodogram'}, ylabel='Variance'>
```



```

[34]: print('p-value of ADF test:')
print(adfuller(df.loc[df.index >= '1964-01-02', 'temp_sa'].dropna())[1])
print('p-value of KPSS test:')
print(kpss(df.loc[df.index >= '1964-01-02', 'temp_sa'].dropna())[1])

```

p-value of ADF test:
0.0
p-value of KPSS test:
0.01

```

[35]: df['temp_sa_d'] = df['temp_sa'].diff()
print('p-value of ADF test:')
print(adfuller(df.loc[df.index >= '1964-01-02', 'temp_sa_d'].dropna())[1])
print('p-value of KPSS test:')

```



```
print(kpss(df.loc[df.index >= '1964-01-02', 'temp_sa_d'].dropna())[1])
```

p-value of ADF test:

0.0

p-value of KPSS test:

0.1

```
[36]: cmap_data = plt.cm.Paired
      cmap_cv = plt.cm.coolwarm


def plot_cv_indices(cv, X, y, group, ax, n_splits, lw=20):
    """Create a sample plot for indices of a cross-validation object."""
    # Generate the training/testing visualizations for each CV split
    for ii, (tr, tt) in enumerate(cv.split(X=X, y=y, groups=group)):
        # Fill in indices with the training/test groups
        indices = np.array([np.nan] * len(X))
        indices[tt] = 1
        indices[tr] = 0
        # Visualize the results
        ax.scatter(
            range(len(indices)),
            [ii + 0.5] * len(indices),
            c=indices,
            marker="_",
            s=50,
            lw=lw,
            cmap=cmap_cv,
            vmin=-0.2,
            vmax=1.2,
        )
        # Formatting
        yticklabels = list(range(n_splits))
        ax.set(
            yticks=np.arange(n_splits) + 0.5,
            yticklabels=yticklabels,
            xlabel="Time series sample index",
            ylabel="CV iteration",
            ylim=[n_splits + 0.2, -0.2],
            xlim=[0, 100],
        )
        ax.set_title("{}".format(type(cv).__name__), fontsize=15)
    return ax

fig, ax = plt.subplots()
cv = TimeSeriesSplit(5, gap=1)
```

```

rng = np.random.RandomState(2024)
X = rng.randn(100, 10)
percentiles_classes = [0.33, 0.33, 0.34]
y = np.hstack(
    [[ii] * int(100 * perc) for ii, perc in enumerate(percentiles_classes)]
)
group_prior = rng.dirichlet([2] * 10)
groups = np.repeat(np.arange(10), rng.multinomial(100, group_prior))
plot_cv_indices(cv, X, y, groups, ax, 5)
ax.legend(
    [Patch(color=cmap_cv(0.8)), Patch(color=cmap_cv(0.02))],
    ["Testing set", "Training set"],
    loc=(1.02, 0.8),
)
# Make the legend fit
plt.tight_layout()
fig.subplots_adjust(right=0.7)

```



```

[37]: # Time series split
tscv = TimeSeriesSplit(n_splits=10, gap=15)
splits = list(tscv.split(df.ld_Close.dropna()))

```

```

[38]: trues_raw = []
preds_raw = []
results_ols_m = []
for i in tqdm.tqdm_notebook(range(len(splits))):
    mod = sm.OLS(
        df.ld_Close.dropna().to_numpy()[splits[i][0]],
        sm.add_constant(
            df.ld_Close.dropna().to_numpy()[splits[i][0]]

```

```

        )[:, [0]]
    )
    res = mod.fit(dispatch=False)
    pred = res.predict(
        sm.add_constant(
            df.ld_Close.dropna().to_numpy()[splits[i][1]]
        )[:, [0]]
    )
    preds_raw.append(pred)
    trues_raw.append(df['ld_Close'].dropna().to_numpy()[splits[i][1]])
    results_ols_m.append(res)

trues = np.concatenate(trues_raw)
preds = np.concatenate(preds_raw)
reg_mean_absolute_error = mean_absolute_error(trues, preds)

linreg_mean_mae = []
for i in range(len(trues_raw)):
    linreg_mean_mae.append(mean_absolute_error(trues_raw[i], preds_raw[i]))

print('MAE, regress on constant alone: ' + str(reg_mean_absolute_error))
del mod, res, pred

```

0%| | 0/10 [00:00<?, ?it/s]

MAE, regress on constant alone: 0.006796108942614544

```

[41]: # ARMA lag order selection using just one fold.
      # This code may run for a minute or two.
      # Feel free to grab a coffee before continuing!

min_ar_ma = [2,6] # Minimum (p, q)
max_ar_ma = [4,8] # Maximum (p, q)

# Note: according to the AIC criteria, identical AR and MA lags are found if
#       the maximum and minimum bounds are:
#min_ar_ma = [1,1] # Minimum (p, q)
#max_ar_ma = [8,8] # Maximum (p, q)

aic_pd = pd.DataFrame(
    np.empty((max_ar_ma[0]+1-min_ar_ma[0],
              max_ar_ma[1]+1-min_ar_ma[1])),
    dtype=float),
    index=list(range(max_ar_ma[0]+1-min_ar_ma[0])),
    columns=list(range(max_ar_ma[1]+1-min_ar_ma[1]))
)

```

```

bic_pd = pd.DataFrame(
    np.empty((max_ar_ma[0]+1-min_ar_ma[0],
              max_ar_ma[1]+1-min_ar_ma[1])),
    dtype=float,
    index=list(range(max_ar_ma[0]+1-min_ar_ma[0])),
    columns=list(range(max_ar_ma[1]+1-min_ar_ma[1]))
)

for p in tqdm.tqdm_notebook(range(
    min_ar_ma[0], max_ar_ma[0]+1), position=1, desc='p'):
    for q in range(min_ar_ma[1], max_ar_ma[1]+1):
        if p == 0 and q == 0:
            aic_pd.loc[p, q] = np.nan
            bic_pd.loc[p, q] = np.nan
            continue
        # Estimate the model with no missing datapoints
        mod = sm.tsa.statespace.SARIMAX(
            df['ld_Close'].dropna().iloc[splits[-1][0]],
            order=(p, 0, q),
            trend='c',
            enforce_invertibility=False
        )
        try:
            res = mod.fit(dispatch=False)
            aic_pd.loc[p, q] = res.aic
            bic_pd.loc[p, q] = res.bic
        except:
            aic_pd.loc[p, q] = np.nan
            bic_pd.loc[p, q] = np.nan

print('AIC: optimal AR order: ' +
      str(aic_pd.min(axis=1).idxmin()) +
      ', optimal MA order: ' +
      str(aic_pd.min().idxmin()))
print('BIC: optimal AR order: ' +
      str(bic_pd.min(axis=1).idxmin()) +
      ', optimal MA order: ' +
      str(bic_pd.min().idxmin()))

```

p: 0% | 0/3 [00:00<?, ?it/s]

AIC: optimal AR order: 2, optimal MA order: 8

BIC: optimal AR order: 2, optimal MA order: 7

```

[42]: trues = []
      preds = []
      results = []
      for i in tqdm.tqdm_notebook(range(len(splits))):

```

```

mod = sm.tsa.statespace.SARIMAX(
    df['ld_Close'].dropna().to_numpy()[splits[i][0]],
    order=(2, 0, 7),
    trend='c',
    enforce_invertibility=False
)
res = mod.fit(dispatch=False)
pred = res.predict(
    data=df['ld_Close'].dropna().to_numpy(),
    start=splits[i][1][0],
    end=splits[i][1][-1]
)
preds.append(pred)
trues.append(df['ld_Close'].dropna().to_numpy()[splits[i][1]])
results.append(res)

trues = np.concatenate(trues)
preds = np.concatenate(preds)

arma_absolute_error = mean_absolute_error(trues, preds)
print('ARMA(2,7) MAE: ' + str(arma_absolute_error))

```

0%| | 0/10 [00:00<?, ?it/s]

ARMA(2,7) MAE: 0.006796211036333233

```

[46]: trues_rf = []
preds_rf = []
X_orig = df[['ld_Close']].dropna()
features = []
for i in range(1,3):
    features.append(
        df[['ld_Close']].dropna().shift(i).rename(
            {'ld_Close': 'ld_Close_'+str(i)}, axis=1
        )
    )
X = pd.concat(features + [X_orig], axis=1)
y = deepcopy(X_orig[['ld_Close']])
X.drop('ld_Close',axis=1,inplace=True)
for i in tqdm.tqdm_notebook(range(len(splits))):
    regr = RandomForestRegressor(criterion="mae",
                                n_estimators=10,
                                max_depth=2,
                                random_state=2024)

    train_idx = splits[i][0][2:]
    res = regr.fit(X.iloc[train_idx],y.iloc[train_idx])
    pred = regr.predict(X.iloc[splits[i][1]])
    preds_rf.append(pred)

```

```

    trues_rf.append(
        df[['ld_Close']].dropna().ld_Close.to_numpy()[splits[i][1]]
    )

trues_rf = np.concatenate(trues_rf)
preds_rf = np.concatenate(preds_rf)

print(str(mean_absolute_error(trues_rf, preds_rf)))

rf_lags_absolute_error = mean_absolute_error(trues_rf, preds_rf)
print('Random forest with lagged prices MAE: ' +
      str(rf_lags_absolute_error))
del regr, res, pred

```

0%| | 0/10 [00:00<?, ?it/s]

0.006781133765126767

Random forest with lagged prices MAE: 0.006781133765126767

[]: